

Delphi

A 0DTE iron condor agent on Alpaca paper trading

Short strikes are set by conformal risk control. The trade condition compares the certified cost of that interval with the market's price of the same interval.

Solo entry · MIT licence · github.com/jannissio/delphi-alpaca-agent

Competition paper account **PA31SEVJV9P9**, funded \$100,000, options level 3, traded only by the submitted agent from 2026-09-03.



Reverse of a silver tridrachm of Delphi, c. 480 BC: four coffers, a dolphin in each. Two maxims stood at the temple, nothing in excess and know thyself. They correspond to the loss caps and to the claims this deck does not make.

Premise: what the evidence supports

Over three trading sessions the profit and loss of an options strategy cannot be distinguished from noise. The agent is therefore built to produce evidence that needs no statistical power.

Traded index-option alphas have been indistinguishable from zero for fifteen years, and same-day premium selling is roughly zero after costs. The one public ODTE rule study lost every positive net Sharpe ratio when its author corrected a hundredfold transaction-cost error in August 2026. [1][2][19]

No edge is claimed, and the corrected result is cited against us.

The half of the index-option premium that can be measured is earned overnight: delta-hedged S&P 500 returns to the buyer average about -1% per day from close to open against $+0.3\%$ from open to close. [3]

Delphi never holds overnight, so it forgoes that half by construction. Pre-registered in advance, not explained afterwards.

With three observations the Probabilistic Sharpe Ratio cannot reach 95 % at any performance level, and three perfect sessions at a credit of 0.20 of the wing cannot push an anytime-valid p-value below 0.51. [4][15]

No Sharpe ratio, win rate or annualised return is reported anywhere in the repository.

Reported instead: a bound on the expected payout of every trade, 31 logged risk gates, and a per-session ledger that `scripts/reproduce.py` regenerates and marks MATCH or MISMATCH. [R11]

[1] Dew-Becker & Giglio 2025. [2] Vilkov 2026 and its KNOWN-ISSUES, 2026-08. [3] Muravyev & Ni 2020 JFE. [4] Bailey & Lopez de Prado 2014. [15] Waudby-Smith & Ramdas 2023. [19] Almeida, Freire & Hizmeri 2025. In the repository: docs/THEORY.md sections 8 and 9, docs/evidence.md, README.md 'Why this design'. [R11] scripts/reproduce.py.

Setting the strikes: conformal risk control

The short strikes sit at the smallest interval whose expected payout to the buyer is certified, in finite samples, to be at most 10 % of the wing.

1 Score

For every past session, how far SPY actually moved divided by the move the options implied. The same unit in history and live.

2 Loss


A condor pays its buyer $\min((r - k)+, w) / w$, a loss between 0 and 1 that falls as the interval widens. [R1, Lemma 1]

3 Radius

From the trailing 250 scores, the smallest radius whose finite-sample corrected payout ratio is at most $\beta^* = 0.10$. [5][R1, Theorem 3]

4 Level

After every session, traded or not, an online level may tighten the radius and never widen it. [6][R1, Theorem 4]




[5] Angelopoulos, Bates, Fisch, Lei & Schuster, ICLR 2024. [6] Feldman, Ringel, Bates & Romano, TMLR 2023. In the repository: [R1] docs/THEORY.md Lemma 1, Theorems 3 and 4; docs/risk_curve.svg, regenerated by scripts/risk_curve.py; [R10] research/G_conformal_condor.md.

The trade condition: market price against certified cost

Sell only when the market pays at least the certified cost of the interval plus the modelled cost of trading it: $\text{credit} / \text{wing} \geq 0.10 + 0.05$.

$$\text{interval} = \text{spot} \pm \text{certified radius} \times \text{implied move}$$



The horizontal axis is the credit as a fraction of the wing. The market's price of the interval, $\text{credit} / \text{wing}$, is read from the quote at the expected fill; the mid is logged alongside.

Back-fill, 618 sessions. History carries no option prices, so only the certified side can be replayed. [R2]

realised payout ratio vs 0.10	all	2024	2025	2026
conformal radius	0.079	0.070	0.079	0.090
fixed 1.10x rule	0.104	0.119	0.113	0.073

[7] Breeden & Litzenberger 1978. In the repository: [R1] docs/THEORY.md Lemma 2 and gate 31; [R2] docs/conformal_backfill.md, 618 calibrated sessions 2024-01-02 to 2026-09-01; [R5] docs/CONFIG_CHANGES.md, credit read at the expected fill from 2026-09-02.

Why credit / wing is the right comparison

For a vertical spread the credit divided by the width is the risk-neutral survival probability averaged across the wing band. It is the market's price of exactly the event the certificate bounds. [7][R1, Lemma 2]

What the condition buys

One line, no free parameter. Every trade that passes is certified, in expectation and under the stated assumptions, not to lose after the modelled round-trip cost of four legs. A trade on the boundary expects its cost and no more.

What it does not buy

The 0.05 is a cost margin, not a profit target, and the certificate is marginal while the condition itself selects sessions. If the market never pays, the agent never trades and the ledger records a closed gate.

Categories from the model, numbers from code

The language model returns enumerations and may only reduce risk. Deterministic code sets every number, and a test proves the order is byte-identical across different model outputs.

Inputs	DATA	Alpaca option chain snapshots, IEX bars and news; Cboe VIX, VIX3M and VIX1D.
Regime	LLM	Enums only: VOL_REGIME, TREND, EVENT_RISK, strategy family, veto. Tickers and dates masked [8]. Three calls must agree, which raises the NO_TRADE rate rather than accuracy [9].
Strategy	CODE	Certified radius to short strikes and wings; credit / wing read from the quote.
Sizing	CODE	2 % of capital per session, 6 % across the campaign; a logistic regime model trained on 9,230 sessions may only shrink the size, never raise it.
Gates	CODE	31 checks, evaluated on every candidate order and logged pass by pass; gate 31 is the trade condition.
Critic	LLM	Sees the finished, gate-validated order and may say PASS, REDUCE or BLOCK. It can never enlarge.
Execution	CODE	Multi-leg limit ladder priced in ticks and re-quoted at every rung; flat by 15:15 ET.
Record	CODE	Append-only audit log with git and configuration hashes, from which the report, the public ledger and the determinism and leakage audits are computed.

Sampling is pinned on temperature, top_p, top_k and seed, and per-field decision entropy is logged in bits [10]. Determinism is an assurance property, not a correctness one: 1,000 temperature-zero completions of one prompt have been measured to give 80 distinct outputs [11]. Open-weight models through Featherless.ai — DeepSeek-V3.2 for regime and critic, Qwen3-30B-A3B for the journal.

[8] Glasserman & Lin 2023. [9] Kim 2026. [10] Koviazin, Mudarisov, Polyachenko & State, IC-AIF 2026. [11] Thinking Machines Lab 2025. In the repository: README.md 'Architecture', docs/WRITEUP.md 'AI logic', [R7] docs/report_2026-09-02.md (LLM section).

Risk gates and execution

Every candidate order passes 31 gates before it is sent, and orders are priced in ticks and re-quoted at every rung.

Thirty-one gates [R4][12]

- 30 controls derived from SEC Rule 15c3-5, MiFID II RTS 6, FINRA Notice 15-09 and the SEC's Knight Capital order, plus gate 31, the trade condition.
- Session loss capped at 2 % of capital with a continuous taper toward a 6 % campaign cap; \$1,000 of maximum loss and 5 contracts per order; defined risk only, so every short leg has a bought wing in the same package.
- Entry windows by weekday; Friday is a logged NO_TRADE day because of the payrolls release at 08:30 ET; no entry after 14:00; flat by 15:15 ET on liquidity grounds.
- Kill switch checked every cycle with cancel-all under five seconds, a daily loss kill on an independent code path, and order and position reconciliation every cycle. A mismatch halts new risk, and did so on the pilot day. [R7]

Execution [13][R5]

- Same-day quoted spreads are one or two ticks wide, so collars and ladders are expressed in ticks and never as a percentage of the mid. No market orders, ever.
- Alpaca paper fills only marketable limit orders. On the pilot day three rungs fixed at 0.49, 0.48 and 0.47 expired unfilled while the package's natural price fell from 0.47 to 0.39 in sixty seconds.
- Since then every rung is priced from quotes pulled at send time and reaches the live natural within sixty seconds, and no rung may go below the credit floor the gates approved.
- Every parameter changed after the first live cycle is listed with its date, its evidence and its effect. [R5]

On Alpaca paper enforced at start-up	Trading API (alpaca-py) chains, multi-leg orders, positions, IEX bars for the calibration scores	Alpaca CLI independent monitoring and reconciliation path, JSON to disk	Alpaca MCP server operator inspection of account and chains
--	--	---	---

[12] SEC Rule 15c3-5; MiFID II RTS 6; FINRA Notice 15-09; SEC order in re Knight Capital, 2013. [13] Fu, Li & Musto, SEC DERA. In the repository: [R4] config/risk_limits.yaml and agent/gates/engine.py; [R5] docs/CONFIG_CHANGES.md; [R7] docs/report_2026-09-02.md.

Live operation and results

The evaluation object is the per-session ledger of the certified cost against the market's price. It is written even when nothing trades.

[dashboard — to be inserted]

In the video this frame carries the screen capture: the dashboard, one full cycle in the audit log, the CLI status call and one MCP call.

Session activity

cycles, gate evaluations, orders, package fills, time to flat

to be filled Friday

P versus Q

credit / wing at the fill against the 0.15 condition, certified payout, gap, decision, and what the fixed rule would have done

to be filled Friday

Where it said no

gate rejections per gate and the deduplicated NO_TRADE reasons

to be filled Friday

Against a null

profit and loss versus SPY over the same hours, and its percentile in the pre-registered random-entry Monte Carlo null: median +259, P05 –503 dollars

to be filled Friday

The 2026-09-02 pilot session — one contract, the fixed-strike rule, +\$14 — ran on a separate development account and is labelled as such in the repository. [R7]

In the repository: [R8] docs/ledger.json and docs/index.html; [R3] docs/evidence.md; [R7] docs/report_2026-09-02.md, the pilot session on the development account; [R11] scripts/reproduce.py. Monte Carlo null: README.md 'What was trained on history'.

Scope of the claim, and next steps

What is machine-checked, what is assumed, what is not claimed, and what more data would change.

Proved

13 theorems in Lean 4 with Mathlib, no sorry, on the three standard axioms.

The condor loss lies in $[0, 1]$ and is non-increasing in the radius; the payoff is affine in the loss; and if the average calibration loss is at most β and the credit is at least $(\beta + \text{margin}) \times \text{wing}$, the average payoff is at least $\text{margin} \times \text{wing}$.

Option pricing has been formalised before and we cite it. A machine-checked risk-control lemma for a trading decision rule we did not find. [R6][18]

Assumed

Exchangeability of the calibration scores.

The certificate is marginal, while the trade condition selects sessions through a quantity that depends on the certified radius. Validity conditional on the gate opening does not follow; the exact repair needs a credit history the basic data plan does not have. [17]

The e-process null is conditional and therefore stronger than the theorem. [R1, section 6 and Theorem 3, remark v]

Not claimed

No Sharpe ratio, win rate or annualised return. The Sharpe-maximising manipulation is precisely our instrument, so the number would be uninformative even if it were computable. [4][14]

No intraday variance risk premium. The measurable half is overnight and is forgone by construction. [3][R1, section 8]

Not the first conformal trading decision: that prior art exists and is cited. [16]

The online level currently has no effect on the interval and is a one-sided safety valve.

Next

A 1DTE overnight condor entered near 15:00 ET, calibrated on the close-to-close horizon where the regime model has fifteen validation blocks.

An option-price history, so the market side of the ledger can be back-tested rather than only logged.

The 10:30 series extended to 2018, so the model comparison gains power.

Mondrian conditioning on the gate event; five votes instead of three.

[4] Bailey & Lopez de Prado 2014. [14] Goetzmann, Ingersoll, Spiegel & Welch 2007. [16] Lekeufack et al., ICRA 2024; Ryan 2026. [17] Jin & Ren 2025; Gibbs, Cherian & Candes 2025; Xu, Guo & Wei 2025. [18] Ushakov & Berdinsky 2026; Coelho 2026. In the repository: [R6] lean/Delphi/Condor.lean and lean/README.md; [R1] docs/THEORY.md sections 6 and 8; [R12] research/STATE_OF_THE_ART.md section 14.

Backup: the certified interval replayed through history

618 calibrated sessions, 2024-01-02 to 2026-09-01. The realised payout ratio stays below the certificate in every year; the fixed-strike rule drifts with whatever the year's realised-to-implied ratio happens to be.

	sessions	realised payout ratio	target	mean radius	fixed rule 1.10x	coverage
All calibrated sessions	618	0.0792	0.10	0.809	0.1039	0.806
2024	202	0.0703	0.10	0.972	0.1188	0.832
2025	249	0.0791	0.10	0.795	0.1126	0.799
2026	167	0.0901	0.10	0.630	0.0729	0.784
Since 2024-12-30 (research/G sample)	418	0.0831	0.10	0.730	0.0962	0.792

State after the back-fill

Online risk level $\beta = 0.1642$ and coverage level $\alpha = 0.2180$ after 618 sessions, window 250 scores, clip [0.35, 1.60], step 0.005.

Because the realised payout ratio has run below β^* , the online level has drifted above it: today the online layer is a one-sided safety valve with no effect on the interval, which we report rather than dress up.

Two limits of this evidence

The whole calibration window lies inside the 2024 to 2026 regime, so it is not a test across regimes.

The coverage track — controlling how often the close leaves the interval — is kept only as the logged counterfactual: most breaches are partial, and the break-even inside probability is 75.7 %, not 83 %.

Repo: [R2] docs/conformal_backfill.md, replayed 2026-09-02; [R10] research/G_conformal_condor.md. Score = $| \text{close} / \text{price at 10:30} - 1 |$ divided by the VIX-implied expected absolute daily move; wing 0.50 % of spot. No option prices exist in this history, so the market side of the ledger cannot be back-tested.

Backup: the regime model and its negative results

A twelve-coefficient logistic regression that can only shrink the position. Two 2026 tabular foundation models were run under the same protocol and did not beat it where anything can be measured.

What it is

9,230 sessions assembled from Cboe (S&P 500 since 1975, VIX since 1990) and Alpaca IEX (SPY open and close from 2020-07, 30-minute bars from 2024-01). A standardised logistic regression on 1,533 SPY sessions maps regime features known at the morning entry to the probability that the session closes inside the short strikes.

What lost

XGBoost, LightGBM, a logit and XGBoost average, TabPFN v2 and TabICL v2, run under the identical expanding-window protocol: seven families across three horizons, 21 configurations. The foundation models win only where nothing can be measured and lose, with bootstrap intervals excluding zero, where it can.

What it may do

Set the size multiplier to 1, 0.5 or 0 — full at or above the historical tercile threshold, half below it, zero in the bottom decile. It can never increase a position.

What the back-test cannot settle

It assumes a constant credit of 17 % of the wing, which invalidates any conclusion that changes the interval width. The live agent reads the credit off the quote instead.

What it cannot do

On the exact 10:30-to-close horizon the logit is worse than a constant, with two year-blocks and no inference possible. On the close-to-close horizon, with fifteen blocks, it is the best single model. Both results are reported.

The null it supplies

The same history gives the random-entry Monte Carlo null against which campaign profit and loss is reported as a percentile rather than as alpha: median +259 dollars, 5th percentile –503 dollars over 2.5 sessions at two contracts per session.

Repo: [R9] docs/regime_model_report.md and research/H_tabular_ml_small_data.md, 21 configurations; README.md 'What was trained on history'. Data: scripts/history_data.py.

Backup: what three sessions can and cannot show

Two test supermartingales are computed from the ledger and reported, never used to halt. The ceiling on the second one is arithmetic, and it is the reason no Sharpe ratio appears anywhere.

Risk process — evidence against the certificate

Null: the expected payout ratio given the past is at most $\beta^* = 0.10$. A p-value below 0.05 would reject the certificate. It does not.

	sessions	mean payout ratio	anytime p
All	619	0.079	0.376
2024	202	0.070	0.376
2025	249	0.079	0.143
2026	168	0.090	1.000

Running maximum of the wealth process 2.659 over 619 sessions. A conformal test martingale on clean data has been observed to fire in 135 of 135 runs at the 5 % level, so this process is reported and never used to halt.

Profit process — and its ceiling

T maximal wins at a credit of g times the wing cannot produce more than $(1/(1-g))^T$. The smallest attainable anytime-valid p-value is its reciprocal.

credit/wing	T = 1	T = 3	T = 5	T = 10	sessions for p <= 0.05
0.15	1.18	1.63	2.25	5.08	19
0.20	1.25	1.95	3.05	9.31	14
0.25	1.33	2.37	4.21	17.76	11

At a credit of 0.20 of the wing, three perfect sessions reach 1.95, so the p-value cannot fall below 0.51. Fourteen consecutive perfect packages would be needed for 0.05. Evidence accrues per package, so two smaller packages a session raise the three-session ceiling to 3.81 at an unchanged risk budget.

Repo: [R3] docs/evidence.md, regenerated by scripts/evidence.py; [R1] docs/THEORY.md section 9. [15] Waudby-Smith & Ramdas 2023; Ramdas, Grunwald, Vovk & Shafer 2023. Bets pre-registered at $\lambda = 1.0$ and $\eta = 1.0$.

Sources: literature (1 of 2)

Not part of the presentation itself. The numbers in square brackets on the slides refer to these entries, so that every claim can be traced back to where it comes from.

- [1] Dew-Becker, I. & Giglio, S. (2025). The Decline of the Variance Risk Premium: Evidence from Traded and Synthetic Options. Federal Reserve Bank of Chicago WP 2025-17.
- [2] Vilkov, G. (2026). ODTE Trading Rules (SPXW, n = 1,319), with the author's KNOWN-ISSUES record of the August 2026 transaction-cost correction, github.com/vilkovgr/0dte-strategies.
- [3] Muravyev, D. & Ni, X. (2020). Why do option returns change sign from day to night? *Journal of Financial Economics* 136(1):219-238.
- [4] Bailey, D. & Lopez de Prado, M. (2014). The Deflated Sharpe Ratio. *Journal of Portfolio Management* 40(5).
- [5] Angelopoulos, A. N., Bates, S., Fisch, A., Lei, L. & Schuster, T. (2024). Conformal Risk Control. ICLR 2024; arXiv:2208.02814.
- [6] Feldman, S., Ringel, L., Bates, S. & Romano, Y. (2023). Achieving Risk Control in Online Learning Settings. TMLR; arXiv:2205.09095.
- [7] Breeden, D. T. & Litzenberger, R. H. (1978). Prices of State-Contingent Claims Implicit in Option Prices. *Journal of Business* 51(4).
- [8] Glasserman, P. & Lin, C. (2024). Assessing Look-Ahead Bias in Stock Return Predictions Generated by GPT Sentiment Analysis. *Journal of Financial Data Science* 6(1).
- [9] Kim (2026) and Bahuguna (2026) on majority voting across size-3 model subsets. Full entries in `research/STATE_OF_THE_ART.md`, section 14.
- [10] Koviazin, Mudarisov, Polyachenko & State (2026). IC-AIF 2026, DOI 10.1145/3800973.3801029.

Sources: literature (2 of 2)

Regulation is taken as a design template for the gates; no compliance is claimed. Entries [16] to [18] are prior art that this work does not claim as its own.

- [11] Thinking Machines Lab (2025), on inference kernels that are not batch-invariant.
- [12] SEC (2010). Rule 15c3-5, Release 34-63241. ESMA / EU (2017). MiFID II RTS 6, Delegated Regulation 2017/589, Articles 12, 15 and 17. FINRA (2015). Regulatory Notice 15-09. SEC (2013). In the Matter of Knight Capital Americas LLC, Release 34-70694. Taken as a design template; no compliance is claimed.
- [13] Fu, L., Li, S., Musto, D. K. & Pearson, N. D. (2025). Hope at a Reasonable Price: Customer Use of Limit Orders in the ODTE Market. SEC DERA Working Paper, 16 March 2025.
- [14] Goetzmann, W., Ingersoll, J., Spiegel, M. & Welch, I. (2007). Portfolio Performance Manipulation and Manipulation-proof Performance Measures. RFS 20(5):1503-1546.
- [15] Waudby-Smith, I. & Ramdas, A. (2023). Estimating means of bounded random variables by betting. JRSS-B 86(1). Ramdas, A., Grunwald, P., Vovk, V. & Shafer, G. (2023). Game-theoretic statistics and safe anytime-valid inference. Statistical Science 38(4).
- [16] Lekeufack, J., Angelopoulos, A. N., Bajcsy, A., Jordan, M. I. & Malik, J. (2024). Conformal Decision Theory. ICRA 2024. Ryan (2026). Conformal Kelly. arXiv:2608.01494. Prior art we do not claim.
- [17] Jin, Y. & Ren, Z. (2025). JRSS-B 87(4). Gibbs, I., Cherian, J. J. & Candes, E. J. (2025). JRSS-B 87(4). Xu, Guo & Wei (2025). Selective Conformal Risk Control. arXiv:2512.12844.
- [18] Ushakov & Berdinsky (2026). arXiv:2608.19223. Coelho (2026). arXiv:2606.01356. Echenim, Guiol & Peltier (2018). arXiv:1807.09873. Prior formal verification in finance.
- [19] Almeida, C., Freire, G. & Hizmeri, R. (2025). ODTE Asset Pricing. SSRN 4701401. Beckmeyer, H., Branger, N. & Gayda, L. (2023). Retail Traders Love ODTE Options... But Should They? SSRN 4404704.

Sources: this repository

Every figure on the slides is regenerated by a script in the repository; `scripts/reproduce.py` recomputes all of them and prints MATCH or MISMATCH against the documents.

- | | | | |
|------|---|-------|---|
| [R1] | <code>docs/THEORY.md</code> — definitions, Lemmas 1, 2 and 5, Theorems 3 and 4, the assumption list in section 6, the pre-registered give-up of the overnight premium in section 8, the e-processes in section 9. | [R7] | <code>docs/report_2026-09-02.md</code> — the pilot session, on the development account PA314NYH4H7G. |
| [R2] | <code>docs/conformal_backfill.md</code> — 618 calibrated sessions replayed through both online tracks. | [R8] | <code>docs/ledger.json</code> and <code>docs/index.html</code> — the published per-session ledger. |
| [R3] | <code>docs/evidence.md</code> — the two e-processes and the evidence ceiling. | [R9] | <code>docs/regime_model_report.md</code> and <code>research/H_tabular_ml_small_data.md</code> — 21 model configurations and their negative results. |
| [R4] | <code>config/risk_limits.yaml</code> and <code>agent/gates/engine.py</code> — the 31 gates as code, immutable at runtime. | [R10] | <code>research/G_conformal_condor.md</code> — the study behind the conformal condor, including the refutation of the width-changing variant. |
| [R5] | <code>docs/CONFIG_CHANGES.md</code> — every pre-registered parameter changed after the first live cycle, with date, evidence and effect. | [R11] | <code>scripts/reproduce.py</code> — regenerates every number quoted in the documents and prints MATCH or MISMATCH. |
| [R6] | <code>lean/Delphi/Condor.lean</code> and <code>lean/README.md</code> — 13 theorems, no sorry, and the narrow statement of what is new. | [R12] | <code>research/STATE_OF_THE_ART.md</code> — the synthesis, and section 14, the independent literature audit of 2026-09-02. |